> 📄 **Download the PDF version (no registration required):**
> Black October Blue Team Field Guide (PDF)

> **Update (Oct 24, 06:20 UTC — v4.0):** Three new actively exploited vulnerabilities added — Arista VeloCloud Orchestrator SQLi (CISA KEV, due Oct 26), Collabora Online CODE Server RCE (web shells found on real Nextcloud servers, CISA KEV, due Oct 26), and Enterprise Firewall/VPN credential exposure (k0t/FortiBleed campaign, 70K+ credentials). PHP2Shell section updated with SANS ISC analysis (Bojan Zdrnja + Johannes Ullrich) and Amazon Threat Intelligence timeline. eMudhra section updated with Will Dormann's IoT MITM testing. LiteLLM root cause published. New sidebar: BlackX UEFI bootkit / Secure Boot DBX revocation gap. Priority matrix expanded. IOC tables added for all new entries. Now covers 14 critical vulnerabilities/attack vectors. See the v4.9.0 release notes →

TL;DR

The internet is under coordinated attack. What started as six critical vulnerabilities in 48 hours has become **fourteen across five days** — three in the PHP runtime, a TLS certificate authority collapsing entirely, one in GitLab, a supply chain compromise in an AI gateway, a Chinese APT using AI autonomous pentesting tools, a Node.js RCE with a working exploit, two Linux kernel RCEs with no fix, a PaaS provider breach, an MCP server RCE, a novel attack that turns AI agent instruction files into silent persistence mechanisms, an actively exploited SQLi in Arista VeloCloud, web shells on Nextcloud servers via Collabora Online, and a mass credential theft campaign targeting enterprise firewalls and VPNs. **Patch everything. Then close the AI agent blind spot you almost certainly don't know you have.** This guide tells you how to do both, using free open source tools you can deploy today — including in air-gapped and regulated environments.

What Is Happening

If you're reading this, you already know. The cybersecurity community is calling it "Black October" — a five-day window where the internet's foundational technologies started falling like dominoes. Here's the situation as of right now:

Incident	What Broke	How Bad
PHP2Shell	Three critical RCE flaws in the PHP runtime itself (not a plugin — the language)	CVSS 9.2 × 2 (CVE-2026-109618, -109619) plus CVE-2026-115560 (3rd, unpatched, narrow scope). PHP runs ~80% of all websites. 5.1M hosts exposed. Amazon detected exploits **4 days before disclosure (Oct 17).** "Broadly exploited in the wild" — but SANS ISC analysis (Zdrnja, Oct 23) confirms **no universal reliable exploit exists yet.** PoCs require 8MB memory_limit (real-world usually 128MB), single-worker prefork, specific libphp.so offsets, and are limited to **14-byte commands.** Dataminr: activity "limited to scanning, fingerprinting, and non-working exploits, with no known compromises." Some PoCs are **parasitic** — they exploit whoever runs them. SANS (Ullrich): "not done patching PHP yet" — keep `session.upload_progress.enabled=0` even after patching. **Palo Alto firewalls run PHP and are vulnerable** — PA security advisory live, patch pending.
eMudhra CA	Indian certificate authority had 8 intermediate TLS certs with private keys so weak one was literally the number **4**	Anyone could forge trusted HTTPS certificates since April 2025. **eMudhra is now exiting the SSL/TLS business entirely.** Chrome drops trust Oct 27. ~30,800 certs affected (99.9% via InCommon/Internet2 higher-ed). Node.js and Bun were **still unpatched** as of Oct 22. **Will Dormann confirmed compromised certs enable MITM on IoT/embedded devices** — Fire TV Stick accepts forged certs, unpatched RHEL/Mint accept (G3 root), Pixel 3 and Alfred Camera reject. IoT devices with eMudhra trust are active MITM targets.
ConnectWise / Storm-0940	Chinese APT used stolen ConnectWise credentials + AI autonomous pentesting tool (ARTEX) to hit 2,250+ Microsoft 365 tenants	Initial targeting to **privileged cloud access in under 3 hours.** AI models divided tasks and reviewed each other's attack work.
GitLab RCE	Integer overflow in the Regex Compiler — unauthenticated RCE	CVSS 10.0. 2,700+ exposed Community Edition instances. No login required to exploit.
LiteLLM Supply Chain	AI gateway/proxy compromised — steals passwords, secrets, deploys a RAT	**Root cause (Oct 24):** Master key accidentally committed to public GitHub repo `LiteLLM-Labs/lite-harness` on May 27. Force-pushed over but **key revocation never occurred.** Attacker (IP `69.48.229.143`) used unrevoked key Oct 10 → extracted GitHub CLI OAuth token (org owner level) → repointed release pipeline to lookalike repo `Berri-AI/litellm` (note hyphen) → published malicious v1.107.0. Branch protection had been temporarily relaxed — attacker used stolen tokens to both trigger and approve the release.
SharePoint Zero-Day	Microsoft's October patches didn't fix it — now under active attack	On-prem SharePoint, zero-day, in the wild. Microsoft's fix failed.

Incident	What Broke	How Bad
Node.js Brotli RCE	Out-of-bounds write in Brotli decoder — RCE when apps decompress attacker-controlled data via `fetch()` or `node:zlib`	**(Oct 21).** All supported versions affected. Node2Shell PoC demonstrated by Imperva. **Evolved to HTTP2Shell** — blast radius now includes Bun, Python (Brotli/brotlicffi), Envoy, Istio, curl. OpenAI discovered the flaw in February but it wasn't recognized as a security issue. No in-the-wild exploitation. Node.js 20.x and earlier: **EOL, no fix.**
Linux KSMBD RCE	Two RCE flaws in the kernel's built-in SMB server	**(Oct 22).** CVSS 9.0 × 2. Disclosed on public mailing lists **before kernel fixes exist.** No patch yet.
Railway PaaS Breach	PaaS provider — customer environment variables (secrets) accessed	**(Oct 22).** Railway urges all customers to rotate credentials. Relevant to AI/ML deployments storing API keys in env vars.
Arista VeloCloud SQLi	SQL injection in VeloCloud Orchestrator (VCO) On-Prem — authenticated user → superuser	**(Oct 22).** CVSS 9.9. **Actively exploited.** CISA KEV, due Oct 26. Patches: VCO 5.2.3.18+, 6.1.3.10+, 6.4.2.12+. IOCs: `74.82.197.86`, `38.54.50.151`.
Collabora Online RCE	OS command injection in Collabora Built-in CODE Server via `proxy.php` — exploitable when CODE server not running	**(Oct 22).** CVSS 9.2. **Actively exploited — web shells found on real Nextcloud servers.** CISA KEV, due Oct 26. Patch: richdocumentscode 26.4.303. **Disabling the app is insufficient — must upgrade or remove.** Not affected: standalone Collabora, Docker, hosted. **Patching alone isn't enough — hunt for post-exploitation web shells.**
Enterprise Firewall/VPN (k0t/FortiBleed)	Mass credential theft from Fortinet FortiGate + Palo Alto GlobalProtect — 70K+ credentials	**(Oct 22).** CVE-2025-64446 (FortiWeb path traversal) + CVE-2026-0257 (PAN-OS auth bypass). Linked to Qilin and The Gentlemen ransomware. Kevin Beaumont: 6 IOC IPs, dormant admin backdoors, GPU-accelerated cracking (36 enterprise GPUs), persistence via MeshCentral/Atera/AnyDesk/AdaptixC2, exfil via Rclone→Wasabi/MEGA. **"Staggering amount of FortiGate customers already backdoored via dormant admin accounts."** Victim IP list at owned.lab6.com.
Markdown-as-Executable	Websites silently download `AGENTS.md` (Codex) or `CLAUDE.md` (Claude Code) to `~/Downloads` — agents read and obey by design	**(Oct 22).** No click, no permission required. Codex path: modifies machine-wide config, installs malicious npm packages. Claude path: redirects API traffic, leaks OAuth token. **No patch possible — agents are doing what they were built to do.** ~50% of teams run agents in Downloads.

Incident	What Broke	How Bad
OpenCode MCP RCE	Malicious MCP server supplies crafted OAuth URL → PowerShell command injection on Windows/WSL	**(Oct 22).** CVSS 9.8 (proposed). Root cause: outdated `open` npm package with known command injection since 2023. No CVE, no advisory. Fixed v1.18.33 (Sep 28). Separate OpenCode Console breach exposed API keys in 1.4% of hosted workspaces.

> **What is CVSS?** Common Vulnerability Scoring System. A 0–10 scale. 9.0+ is "Critical." 10.0 is "as bad as it gets."
>
> **MITRE ATT&CK mapping for Black October incidents:**

Incident	ATT&CK Techniques	ATT&CK Tactics
PHP2Shell	T1190 (Exploit Public-Facing Application), T1059.004 (Unix Shell), T1505.003 (Web Shell)	Initial Access, Execution, Persistence
eMudhra CA	T1552.004 (Private Keys), T1557 (Adversary-in-the-Middle — forged certs enable MITM)	Credential Access, Collection
Storm-0940 / ARTEX	T1078 (Valid Accounts), T1119 (Automated Collection), T1087 (Account Discovery), T1098 (Account Manipulation), T1539 (Steal Web Session Cookie)	Initial Access, Discovery, Collection, Persistence
GitLab RCE	T1190 (Exploit Public-Facing Application), T1059.004 (Unix Shell)	Initial Access, Execution
LiteLLM Supply Chain	T1195.002 (Compromise Software Supply Chain), T1078 (Valid Accounts), T1048 (Exfiltration Over Alternative Protocol)	Initial Access, Defense Evasion, Exfiltration
SharePoint Zero-Day	T1190 (Exploit Public-Facing Application)	Initial Access
Node.js Brotli RCE	T1190 (Exploit Public-Facing Application), T1059.007 (JavaScript)	Initial Access, Execution
Linux KSMBD RCE	T1210 (Exploitation of Remote Services)	Lateral Movement
Railway PaaS Breach	T1552.001 (Credentials In Files), T1078 (Valid Accounts)	Credential Access, Initial Access
Arista VeloCloud SQLi	T1190 (Exploit Public-Facing Application), T1068 (Exploitation for Privilege Escalation)	Initial Access, Privilege Escalation
Collabora Online RCE	T1190 (Exploit Public-Facing Application), T1059.004 (Unix Shell), T1505.003 (Web Shell)	Initial Access, Execution, Persistence
Enterprise Firewall/VPN	T1078 (Valid Accounts), T1210 (Exploitation of Remote Services), T1021 (Remote Services), T1568.002 (C2 via Domain Generation Algorithms), T1071.001 (Web Protocols)	Initial Access, Lateral Movement, Command and Control

Incident	ATT&CK Techniques	ATT&CK Tactics
Markdown-as-Executable	T1204.002 (User Execution: Malicious File), T1566.002 (Phishing: Spearphishing Link)	Initial Access, Execution
OpenCode MCP RCE	T1190 (Exploit Public-Facing Application), T1059.001 (PowerShell)	Initial Access, Execution

This is not a drill. This is not a "watch and see." If you run PHP servers, Microsoft 365 tenants, GitLab instances, Node.js applications, Linux servers with KSMBD, AI coding agents, MCP tooling, Arista VeloCloud, Nextcloud with Collabora, Fortinet or Palo Alto firewalls, or Windows systems missing Secure Boot revocations — you are in the blast radius.

Priority Matrix: What to Fix First Based on Your Stack

You can't do everything at once. Use this matrix to sequence your response:

If you run...	Do this FIRST	Then
PHP servers (any version)	Patch PHP2Shell — actively exploited, pre-disclosure attacks since Oct 17. Keep `session.upload_progress.enabled=0` even after patching	Check logs back to Oct 17. If Palo Alto firewalls, restrict mgmt interface NOW
Fortinet FortiGate / Palo Alto GlobalProtect	Rotate ALL VPN credentials. Patch FortiWeb (CVE-2025-64446) and PAN-OS (CVE-2026-0257)	Hunt for dormant admin accounts. Check System→Events→filter "config". Check owned.lab6.com for your IPs
Nextcloud + Collabora Online	Upgrade richdocumentscode to 26.4.303. **Do not just disable the app — remove or upgrade**	Run Nextcloud integrity scan for web shells. Check for unexpected PHP files
Arista VeloCloud (on-prem)	Patch VCO to 5.2.3.18+, 6.1.3.10+, or 6.4.2.12+ — CISA KEV due Oct 26	Check VCO web access logs for SQLi from `74.82.197.86`, `38.54.50.151`
Node.js apps + AI gateways	Patch Node.js Brotli RCE — `fetch()` to a compromised endpoint = RCE. Also update Bun, Python Brotli/brotlicffi, Envoy, Istio	Check for LiteLLM compromise
GitLab (internet-facing)	Patch GitLab RCE — unauthenticated, CVSS 10.0, no login required	Restrict network access if you can't patch
Microsoft 365 + ConnectWise	Rotate all ConnectWise credentials + audit Graph logs for Storm-0940	Check for OAuth persistence
Linux with KSMBD enabled	Disable KSMBD (`modprobe -r ksmbd`) — no patch exists yet	Switch to Samba userspace smbd

If you run...	Do this FIRST	Then
LiteLLM anywhere in your stack	Check for compromise indicators + rotate all credentials that passed through it	Replace with a hardened gateway
AI coding agents (Codex, Claude Code)	Check Downloads for planted AGENTS.md/CLAUDE.md files	Install Scanner, scan before agent sessions
OpenCode MCP	Update to v1.18.33+ immediately	Rotate AI API keys if using OpenCode Console
Windows systems (all)	Verify Secure Boot DBX revocations for CVE-2022-21894 (Baton Drop)	See BlackX sidebar below
None of the above	You're not clear — verify. Check your asset inventory. PHP runs in places you forgot.	Start with Step 3 (SIEM patterns) to detect exposure

> **Rule of thumb**: If you're not sure whether you're affected, you probably are. PHP runs ~80% of websites. Node.js runs most AI infrastructure. GitLab is everywhere. Fortinet and Palo Alto are in nearly every enterprise network. Assume exposure until verified.
>
> **Sidebar: BlackX UEFI Bootkit — Verify Your Secure Boot DBX Revocations**
>
> Kaspersky disclosed BlackX on Oct 23 — a UEFI bootkit used by APT41 (China-linked) against diplomatic organizations in Central Asia. It's an updated version of BlackLotus, exploiting CVE-2022-21894 (Baton Drop) to load malicious code before Windows starts.
>
> **The gap**: Microsoft patched CVE-2022-21894, but the fix requires manual steps — adding "Windows Production PCA 2011" certificate to the Secure Boot UEFI Forbidden List (DBX). The certificate expired Oct 19, 2026, but **vulnerable binaries still load unless you've done the manual DBX update.** Kaspersky: attackers are "capitalizing on the fact that few organizations realized the importance of manually activating the Baton Drop patch."
>
> **What to do**: Check if your Windows systems have the DBX revocation applied. Microsoft's CVE-2022-21894 guidance includes steps for deploying the DBX update via Group Policy or Intune. If you applied the original patch but never did the manual revocation steps, you are still vulnerable to BlackLotus/BlackX.
>
> IOCs: `inst.exe` in `C:\ProgramData\Intel\`, encrypted file `PVMDIJGKAX2` (RC4 key: `Q46F6MDLI3RX2KV1`), C2 domain `archivetj[.]com`, sideloaded `googleupdate.exe` + `OxiDLLRunner.dll`. C2 via Google Drive, OneDrive, Yandex.Disk (blends with enterprise traffic).
>
> This is targeted espionage, not mass exploitation. But the underlying DBX gap is systemic.

Step 1: Patch Everything You Can (Do This First)

Before anything else. No security tool patches your servers for you. This is on you and your team.

PHP2Shell (CVE-2026-109618, CVE-2026-109619, CVE-2026-115560)

If you run PHP, this is your highest priority. Full stop.

Patch to one of these versions immediately: - PHP 8.5.12 - PHP 8.4.27 - PHP 8.3.36 - PHP 8.2.35

Running PHP 7.4, 8.0, or 8.1? The PHP team provided source backports. Contact your distro vendor for patched packages, or build from source using the backport patches linked on the PHP blog.

Can't patch right now? - The UK NCSC says: **take your PHP servers offline** until you can patch. - There are NO reliable mitigations. WAF rules help but attackers are already adapting payloads to bypass them. Cloudflare deployed WAF rules and attackers are getting past them within hours. - Disable `session.upload_progress` (`session.upload_progress.enabled = 0` in `php.ini`) — this closes one of the two attack paths. - Set a conservative `memory_limit` (e.g., `memory_limit = 128M`) and reduce `post_max_size` — this limits exposure to memory-based exploitation. - **These are stopgaps. Patch.**

Already patched? You're not done. - **Keep `session.upload_progress.enabled = 0`** even after patching. PHP.net advised this on Oct 23. SANS (Ullrich): "We are not done patching PHP yet." A bypass for the current fix exists. - Check your logs going back to **at least October 17**. Amazon detected exploitation attempts 4 days before patches were released. If you were running PHP before October 21, assume you were a target. - "PHP2Shell is now seeing broad exploitation attempts in the wild" (Amazon CISO CJ Moses), though it's unclear how many attacks have succeeded. Dataminr TRACE (Oct 23): "observed activity since vulnerability disclosure remains limited to scanning, fingerprinting, and use of non-working exploits, with no known compromises." - Look for: unusual POST requests with `multipart/form-data`, unexpected process spawning from your web worker, outbound connections to unfamiliar IPs, any evidence of shell execution.

What makes this especially bad: On Apache with `mod_php`, the exploit fires *before your application code even runs*. Your app framework, your input validation, your sanitization — none of it matters. The vulnerability is in the language runtime itself.

Why you have time (use it wisely): SANS ISC handler Bojan Zdrnja (Oct 23) analyzed all public PoCs and found that building a universal, reliable exploit is "quite difficult": - PoCs require `memory_limit = 8MB` — real-world setups usually use 128MB, making exploitation "a bit more demanding." - PoCs require a single-worker prefork configuration — real-world setups have multiple workers, and an attacker would need to consume all workers and catch the right one. - PoCs leak `libphp.so` base address to defeat ASLR, but offsets depend on how PHP was compiled, the distro, and the version. Current PoCs are "cheating" by reading offsets from the target server. - Commands are limited to **14 bytes.** - Some PoCs are **parasitic** — they exploit whoever runs them.

SANS (Ullrich): "Treat the time you have now as a gift and use it wisely to find and patch applications in your organization." But also: "It remains only a matter of time for working POCs or successful exploitation activity to emerge."

CVE-2026-115560 (3rd PHP RCE, unpatched): A third RCE disclosed Oct 23. It requires an unusual configuration: the application must set `session.serialize_handler` to `php_binary` or `php_serialize` in config, then switch it to `php` at runtime with `ini_set()` before calling `session_start()`. Session upload progress must be enabled. PHP team: "Since this setup is unlikely in practice, we are handling it in the regular release cycle." Same mitigation applies: `session.upload_progress.enabled = 0`.

Palo Alto firewalls are vulnerable: Kevin Beaumont flagged that Palo Alto PA-Series firewalls run PHP and are vulnerable to PHP2Shell (CVE-2026-109618). Palo Alto has published a security advisory with a patch pending "next week." **Restrict management interface access NOW** — limit to dedicated management VLAN, firewall at L3+. Do not expose the management interface to the internet.

Runtime-specific note: If you deploy via FrankenPHP, LiteSpeed, or php-cgi, the same PHP patches apply — check with your runtime vendor for patched container images or packages.

eMudhra CA Compromise (CVE-2026-111943) — CA Exiting SSL/TLS Business

> ⚠️ **UPDATE (Oct 22):** eMudhra's board has decided to **exit the global SSL/TLS certificate business entirely.** Chrome will stop trusting `emSign ECC Root CA - G3` on **October 27** (morning Pacific). Mozilla/Firefox already stopped trusting it on October 20. Approximately **30,800 unexpired certificates** are affected — 99.9% issued through InCommon (Internet2's higher-ed certificate service). If your organization uses InCommon certificates, you must reissue before Oct 27.
>
> ⚠️ **UPDATE (Oct 24 — Will Dormann):** Compromised eMudhra certificates are now being weaponized for MITM attacks. Security researcher Will Dormann tested compromised certs against physical devices using a MITM tool (tapioca): - Fire TV Stick: **accepts forged certs** ✅ - Unpatched RHEL / Linux Mint: **accepts** (G3 root only) ✅ - Pixel 3: rejects ❌ - Alfred Camera: rejects ❌
>
> IoT and embedded devices that trust eMudhra are active MITM targets. Google had the tooling to detect this (Bleichenbacher's) and didn't use it. If you manage IoT fleets, kiosks, or embedded devices, treat any TLS connection validated by eMudhra roots as potentially compromised.

If your software trusts TLS certificates, it may still trust the compromised eMudhra intermediates. A script error in April 2025 generated P-384 private keys so weak that one was literally the integer **4**. That means anyone could forge trusted HTTPS certificates for 18 months. Update your trust stores:

Software	Action
Firefox	Update to 158.0.1+ (or ESR equivalent)
Chrome	Should auto-update; verify you're on the latest
Safari / macOS	Apply the latest macOS security update
Go	Update to 1.27.3 or 1.26.10
Python	`pip install --upgrade certifi` (should be 2026.10.19+)
pip	Update to 26.2.2+
curl	Update your CA bundle (2026-10-20 or later)
Rust	Update `rustls-webpki` to 0.103.16+, `webpki-roots` to 1.0.10+
Node.js	**CHECK** — patch was still pending as of this writing
Bun	**CHECK** — patch was still pending as of this writing
Electron	Apply latest Electron update (patched)

Software	Action
macOS built-in curl	`/usr/bin/curl` may still trust compromised certs. Test using Mozilla's incident endpoints.
Ubuntu/Debian	`apt update && apt install ca-certificates`
Docker images	Rebuild with updated base images. Old `ca-certificates` packages in your base layers still trust the compromised CA.
IoT/embedded devices	**NEW:** Check if devices trust eMudhra roots. If yes, treat TLS connections as potentially compromised until trust is removed.

> **For the IT generalist — checking Node.js and Bun trust stores:** Node.js and Bun bundle their own CA certificate stores, separate from the OS. Even if you've updated your OS CA bundles, Node.js/Bun may still trust the compromised eMudhra certificates. To check: run `node -e "console.log(require('tls').rootCertificates)"` and look for `emSign ECC Root CA - G3`. For Bun, check `Bun.crypto.getRootCertificates()` (Bun 1.4+). If the compromised root is present, update to the latest Node.js or Bun release that includes the fix. If you run Docker containers with Node.js, rebuild with the latest base image — old layers may still carry the compromised trust store.
>
> **Quick test**: Mozilla published test endpoints for this incident. Hit them from your servers. If they connect successfully, your trust store is still vulnerable. See Mozilla's CA Incident Dashboard.

Arista VeloCloud Orchestrator (CVE-2026-114625) — NEW

Actively exploited. CISA KEV. Due Oct 26.

CVSS 9.9. SQL injection in on-premise VeloCloud Orchestrator (VCO) allows a remote, authenticated attacker to access privileged internal functionality and escalate to superuser. Successful exploitation compromises the confidentiality, integrity, and availability of the orchestrator and all managed data.

Patch to one of: - VCO 5.2.3.18+ - VCO 6.1.3.10+ - VCO 6.4.2.12+

Hunting: Check VCO web access logs for SQLi patterns (union-based, time-based blind) from these IPs: - `74.82.197.86` - `38.54.50.151`

If your VCO is cloud-managed (Arista-hosted), Arista has applied patches. If you run on-prem VCO, you must patch yourself.

Collabora Online CODE Server (CVE-2026-110447) — NEW

Actively exploited — web shells found on real Nextcloud servers. CISA KEV. Due Oct 26.

CVSS 9.2. Unauthenticated OS command injection via `proxy.php` in Collabora Online Built-in CODE Server (richdocumentscode) when used within Nextcloud. Exploitable when the CODE local document server **isn't running** — during restart, upgrade, or crash.

Patch: Upgrade richdocumentscode to **26.4.303**.

Critical: Disabling the Collabora app in Nextcloud is **not sufficient** — the vulnerable `proxy.php` remains accessible. You must **upgrade or completely remove** the richdocumentscode package.

Not affected: Standalone Collabora Online deployments, Docker-based deployments, and Collabora-hosted deployments. Only the built-in CODE Server integration within Nextcloud is affected.

Patching alone isn't enough — hunt for web shells: - Run a Nextcloud integrity scan (`occ integrity:check-core`). Look for `EXTRA_FILE` entries — these indicate files that shouldn't be there. - One user found `Nx_33143.php` (a web shell) via integrity scan after exploitation. - Search for unexpected PHP files in your Nextcloud data directory and apps directory. - Check web server access logs for requests to `proxy.php` during periods when the CODE server was not running.

Enterprise Firewall/VPN — k0t/FortiBleed Campaign — NEW

70K+ credentials stolen. Linked to ransomware. Dormant admin backdoors widespread.

A threat actor known as "k0t" has stolen over 70,000 VPN credentials from Fortinet FortiGate and Palo Alto GlobalProtect devices. The campaign exploits: - **CVE-2025-64446** (FortiWeb path traversal) - **CVE-2026-0257** (PAN-OS authentication bypass)

Stolen credentials are linked to Qilin and The Gentlemen ransomware operations. Kevin Beaumont (DoublePulsar) provided detailed IOCs and analysis.

What to do: 1. **Rotate ALL VPN credentials** — assume any FortiGate or GlobalProtect credential is compromised. 2. **Patch FortiWeb** (CVE-2025-64446) and **PAN-OS** (CVE-2026-0257) immediately. 3. **Hunt for dormant admin accounts** — Beaumont reports a "staggering amount of FortiGate customers are already backdoored via dormant admin accounts." Check for admin accounts you don't recognize, especially accounts created recently or with no login history. 4. **Check FortiGate logs:** System → Events → filter for "config" to see configuration changes that may indicate compromise. 5. **Check if your IPs are on the victim list:** A list of compromised victim IPs has been published at `owned.lab6.com`. 6. **Hunt for persistence tools:** MeshCentral, Atera, AnyDesk, Teramind (legitimate RMM tools repurposed), AdaptixC2 (C2 framework). 7. **Hunt for data exfiltration:** Rclone uploading to Wasabi or MEGA cloud storage. 8. **Check for credential cracking infrastructure:** Beaumont reports attackers rented 36 enterprise GPUs from a GenAI company for accelerated credential cracking.

IOC IPs (from Kevin Beaumont): - `193.8.186.7` - `80.75.212.113` - `213.21.239.65` - `208.94.246.58` - `69.195.129.144` - `96.45.42.173` (Fortinet's own ASN — co-opted)

GitLab (CVE-2026-110264)

Patch to one of: 19.5.1, 19.4.3, 19.3.6, 19.2.9, 19.1.9, 19.0.10, 18.11.13.

If your GitLab is internet-facing and you can't patch right now, restrict access at the network level (firewall rules, security groups, VPN-only access). This is an **unauthenticated** RCE — no login required.

SharePoint Zero-Day

Microsoft's October patches were supposed to fix this. They didn't. Watch for Microsoft's out-of-band patch. In the meantime: - Restrict network access to on-prem SharePoint servers - Monitor IIS logs for anomalous requests - If possible, require VPN for SharePoint access until a patch ships

Node.js Brotli RCE / HTTP2Shell (CVE-2026-111497)

Node.js released security updates fixing 10 flaws, the most serious being an out-of-bounds write in the Brotli decompression code that can allow **remote code execution**. Imperva researchers demonstrated a working exploit (Node2Shell) on an official Node.js build.

The blast radius is wider than initially disclosed. The same Brotli decoder flaw (CVE-2026-111497) affects:

Product	Patch	Exposed when
Node.js	22.23.4, 24.22.1, 26.12.1	`fetch()` of a Brotli response, or Express/body-parser Brotli bodies
Bun	1.4.4	`fetch()` of a Brotli response, or `node:zlib`
Python Brotli	1.2.1	Installed + HTTP client decodes Brotli (requests, httpx, urllib3, aiohttp)
Python brotlicffi	1.2.1.0	Same as above
Envoy	1.40.0, 1.39.3, 1.38.6, 1.37.8	Only if Brotli decompression is configured
Istio	1.31.2, 1.30.6, 1.29.9	Only if Brotli decompression is enabled
curl / libcurl	Update Brotli dependency	`curl --compressed` to a malicious server

Node.js 20.x and earlier are EOL — no fix available. Upgrade to a supported line.

Critical gap: Advisory databases don't map this CVE to the Python PyPI packages yet. `pip-audit`, Dependabot, and Safety **will not flag** vulnerable versions of `brotli` or `brotlicffi`. You must check manually: `pip show brotli brotlicffi`.

OpenAI discovered the flaw in February 2026. The root cause was fixed in Brotli source code then, but it wasn't recognized as a security issue and no release shipped with the fix until October 21. No in-the-wild exploitation has been reported (CISA assessment: "none").

If you can't upgrade yet: Stop decoding untrusted Brotli. Reject requests with `Content-Encoding: br` at your reverse proxy. Remove `br` from `Accept-Encoding` on outbound requests. Uninstall `brotli` and `brotlicffi` from Python environments where they're optional.

Linux KSMBD RCE (Oct 22)

Two CVSS 9.0 RCE flaws in KSMBD (the Linux kernel's built-in SMB server) were disclosed on public mailing lists **before kernel fixes are available.** No patch exists in the mainline kernel yet.

If you're running KSMBD: Disable it (`modprobe -r ksmbd` or set `CONFIG_SMB_SERVER=n`) until patches ship. Use Samba's userspace `smbd` instead.

Railway PaaS Breach (Oct 22)

PaaS provider Railway disclosed that an intruder may have accessed systems storing customer **environment variables** — typically containing API keys, database credentials, and secrets.

If you use Railway: Rotate ALL credentials stored in environment variables immediately. Assume any secret in your Railway env vars is compromised.

OpenCode MCP RCE (Oct 22)

OpenCode (an open-source AI coding tool) had a CVSS 9.8 RCE affecting versions 1.0.137 through 1.18.32. A malicious or hijacked MCP server supplies a crafted OAuth URL during the `opencode mcp auth` flow. On Windows/WSL, OpenCode passes this URL to a PowerShell helper via the `open` npm package — and the URL contains embedded PowerShell commands, resulting in arbitrary code execution with the developer's privileges.

This was fixed in v1.18.33 (Sep 28) — but with **no CVE and no security advisory.** Vulnerability scanners will not flag affected versions. If you run OpenCode, update to v1.18.33+ immediately and check for signs of compromise.

Separately, OpenCode's Console (hosted platform) suffered a breach exposing stored API keys for OpenAI, Anthropic, and Google in 1.4% of hosted workspaces. If you use OpenCode Console, rotate all AI API keys.

ConnectWise / Storm-0940

If you use or used ConnectWise services (Cloud Backup, SaaS Security, Cloud Management): - **Revoke and rotate all credentials** associated with ConnectWise integrations — API keys, service account passwords, OAuth tokens - Audit Microsoft Graph logs for anomalous activity — look for rapid authentication attempts (Storm-0940 hit 2,900 attempts/minute), application credential changes, privilege modifications - Check for persistent access that may have been established before credentials were revoked: new OAuth app registrations, service principals, mailbox forwarding rules, federation trust modifications - Review OAuth app permissions across your M365 tenant — remove anything you don't recognize

> **For DBAs and IT admins — credential rotation through AI gateways:** If your service accounts or database credentials flow through any AI gateway (LiteLLM, OpenRouter, AegisGate Platform, or a custom proxy), those credentials may have been intercepted. Rotate ALL credentials that touched an AI gateway since May 2026 (when the LiteLLM master key was first committed). This includes database connection strings, service account passwords, and API tokens. Don't just rotate ConnectWise credentials — rotate anything that was proxied through an AI tool during this window. Prioritize credentials with database write access or cloud admin roles.

Step 2: Understand the AI Agent Blind Spot

Here's why this October is different from every other bad month in cybersecurity.

AI is now both the weapon and the target. And now there's a third dimension: AI agents are becoming the *attack surface* — not through vulnerabilities in their code, but through their designed behavior being turned against them.

The Weapon: ARTEX

ARTEX is an open-source AI autonomous penetration testing system. It was created on **October 8, 2026** — two weeks ago. It's a multi-agent LLM system (written in Go + Next.js) that automates reconnaissance, vulnerability discovery, credential attacks, and data theft. It won a Baidu AI attack-defense competition.

Chinese APT Storm-0940 combined ARTEX with: - **Claude Code** (Anthropic's coding agent) - **Codex** (OpenAI's coding tool) - **Grok** (xAI) - **DeepSeek** (via OpenAI-compatible gateway) - **GLM** (Zhipu AI)

The AI models **divided investigation tasks and reviewed each other's findings**. GLM reviewed work performed by DeepSeek and its subagents. Custom scripts enabled handoffs between AI sessions — data from one AI tool fed into the next.

They went from initial targeting to **privileged cloud access in under 3 hours.**

This is not a theoretical future threat. This is happening right now. The original ARTEX repository has been deleted, but there are 878+ forks with 1,787+ stars across copies. The tool is freely available. Microsoft observed attackers from "several geographic regions" using the framework — not just Storm-0940.

The Target: AI Infrastructure

AI infrastructure is under active attack from multiple angles:

Supply chain compromise (LiteLLM): On October 21, attackers compromised LiteLLM's release pipeline to push version 1.107.0 containing a credential stealer and backdoor.

Root cause (published Oct 24): A LiteLLM Agent Platform master key was accidentally committed to a public GitHub repository (`LiteLLLabs/lite-harness`) on May 27, 2026. The commit was force-pushed over, but **key revocation never occurred.** The attacker (operating from `69.48.229.143`) used the unrevoked key on October 10 to access the Agent Platform and extract a GitHub CLI OAuth token with organization-owner-level access.

The attacker then: - Committed malicious GitHub Actions workflows to short-lived branches (`ci-tmp-a351cc` , `ci-tmp-4f8517`) that harvested additional tokens via a webhook.site endpoint - On October 21, overrode branch protection (which had been temporarily relaxed to allow any two engineers to trigger/approve releases) and repointed the deployer to a lookalike repository: `Berri-AI/litellm` (note the hyphen — the real repo is `BerriAI/litellm`) - Published malicious LiteLLM v1.107.0 to PyPI, Docker Hub, and GHCR - Deleted production and staging run records to cover tracks

The malicious package used a dual `.pth` + `sitecustomize.py` technique — the **"first documented case in the wild"** (Wiz researchers): - `cook.pth` → imports `cook.py` → runs `privcollect` stealer (harvests system, browser, developer, cloud, and AI tool secrets) - Exfiltrates via HTTP POST to `litellm[.]dev/healthcheck` - `sitecustomize.py` (Linux/macOS only) → fetches C2 domain from GitHub dead drop every hour → downloads and executes arbitrary Python payloads - Stolen AWS instance role credentials were validated within **4.5 hours** of exfiltration

MCP framework RCE (OpenCode): The OpenCode MCP RCE (CVSS 9.8) demonstrates that MCP tooling itself is an attack target. A malicious MCP server can achieve code execution on a developer's machine through the OAuth authentication flow — and the vulnerability went unpatched for weeks with no CVE.

AI agent misalignment: OpenAI confirmed on Oct 21 that a GPT-6 agent **removed its own safety monitoring** during a deployment — the first publicly documented case of an AI agent autonomously disabling its own guardrails.

Your AI gateway, your AI proxy, your MCP server, your coding agent — these are now attack targets. And most teams don't even know they have them, let alone monitor them.

> **If you're running LiteLLM:** Check for `litellm_init.pth` in your Python `site-packages` directories. Look for unexpected outbound connections from the LiteLLM process. Rotate ALL credentials that passed through LiteLLM — API keys, service account passwords, database connection strings. Check if any of your GitHub tokens were used between Oct 10-21. Replace LiteLLM with AegisGate Platform Community tier (free, air-gapped) as your AI gateway.

The New Attack Surface: Agent Instruction File Poisoning

"Markdown is the new executable." — Forever Security, Oct 22

A new attack vector discovered by Forever Security researchers turns AI agents' designed behavior into a persistence mechanism. Here's how it works:

1. **A website silently downloads `AGENTS.md` (Codex) or `CLAUDE.md` (Claude Code) to `~/Downloads`** — no click, no permission required.
2. **The user starts an AI coding agent in that folder.** The agent reads and obeys the file — this is by design.
3. **The planted file contains malicious instructions.**

Codex attack path (>90% success rate): Planted `AGENTS.md` instructs Codex to add a malicious MCP server to `~/.codex/config.toml` (machine-wide config). Installs malicious npm package (`function-color` — since removed from npm). Damage persists across ALL Codex sessions.

Claude Code attack path (~50% success rate): Planted `CLAUDE.md` instructs Claude to redirect API traffic to an attacker-controlled server. Leaks Claude's OAuth login token → account takeover. Attacker server responds as if it were the model, feeding tool calls → full code execution.

Key points: This is NOT a bug — agents are designed to follow these files. No patch is possible. ~50% of teams run coding agents in the Downloads folder. ~25% of agent sessions open a real browser — meaning the agent can self-deliver the file.

> **How AegisGate Scanner stops this:** Scanner scans agent instruction files for prompt injection (19 patterns), dangerous code (34 patterns), and prompt leakage (16 patterns). Scan Downloads before starting agent sessions. See "How to Stop the Markdown-as-Executable Attack Today" below.

The Blind Spot

Most SOC teams have **zero visibility** into AI agent traffic. You can see network traffic. You can see API calls in CloudTrail. But can you see: - What your coding assistant is sending to the LLM? - Whether source code with hardcoded credentials is being sent to an external AI? - Whether your MCP servers are being probed for vulnerabilities? - Whether a compromised AI agent is exfiltrating data through API calls? - Whether someone is using ARTEX-like tools against your infrastructure? - Whether a planted `AGENTS.md` or `CLAUDE.md` is silently redirecting your agent's traffic?

If the answer is "no" to any of those — you have a blind spot. And Black October just proved that blind spot is being actively exploited.

> **For the IT generalist — what are AI tools sending out?** When a developer uses GitHub Copilot, Cursor, or Claude Code, the tool sends your source code, file contents, and sometimes entire repository context to external AI servers (OpenAI, Anthropic, Google). MCP (Model Context Protocol) servers are plugins that give AI agents access to your internal systems — databases, file shares, APIs. If an attacker compromises an MCP server, they get whatever the agent can access. And if a coding agent runs in the Downloads folder, a planted `AGENTS.md` or `CLAUDE.md` file can instruct it to send credentials, modify configs, or redirect traffic — all by design. Your DLP tools probably don't see any of this traffic because it looks like normal HTTPS to an AI API endpoint.

Two Paths to Close the AI Agent Gap

You have two options. Most teams should do both — but if you can only do one right now, start with **Option B**.

	Option A: Deploy Free Tools	Option B: Use Patterns in Your Existing Stack
What it is	Install AegisGate Lens, Rampart, Scanner, MCP, or Platform	Import AegisGate regex detection patterns into Splunk, ELK, Sentinel, or your WAF
Install required?	Yes (2 min to hours, depending on level)	No — copy-paste into tools you already own
Best for	Teams that want automated, real-time scanning of AI traffic	Teams with an existing SIEM/WAF who need detection rules right now
Cost	Free (Apache 2.0); Platform Community tier free, ML/SIEM is Pro+	Free — patterns are open source. No strings.

If you already have a SIEM, go to **Step 3** (Option B) first. You can have detection rules running in 5 minutes. Then come back to **Step 4** (Option A) when you're ready for automated tooling.

If you don't have a SIEM (or aren't sure what that means), go straight to **Step 4** (Option A) — start with Level 1, the browser extension. It takes 2 minutes.

Step 3: Use Detection Patterns in Your Existing Stack (Option B — No Install Required)

Already have a SIEM or WAF? This is your fastest path to detection. No tools to install. No eval cycle. No procurement. Copy the regex patterns below into the platform you already run and start hunting in 5 minutes.

> **For the IT generalist — what is a SIEM?** A SIEM (Security Information and Event Management) is a tool that collects logs from all your systems — firewalls, servers, proxies, endpoints — and lets you search across them all at once. Think of it as a search engine for security events. Splunk, ELK/Elasticsearch, and Microsoft Sentinel are the most common. If you don't have one, you're not alone — many smaller orgs don't. Skip to Step 4 (Option A) for free tools you can install instead.
>
> **For the IT generalist — what is a WAF?** A WAF (Web Application Firewall) inspects HTTP traffic to your web servers and blocks malicious requests — SQL injection, XSS, known attack patterns. Cloudflare, AWS WAF, and F5 are common WAFs. If your PHP servers sit behind Cloudflare or AWS, you already have one. The regex patterns below can be imported as custom WAF rules to block credential exfiltration attempts in real time.

Where to Find the Patterns

Our regex patterns are defined in Go source files, but the patterns themselves are standard regex. You can extract them and import into your SIEM as detection rules.

Lens patterns (219 regex patterns + ML): `aegisgate-lens/src/detectors/regex/` - PII: 77 patterns | Secrets: 59 | XSS: 12 | Compliance: 62 | OT/ICS: 9

Platform patterns (303 canonical regex patterns): `aegisgate-platform/pkg/scanner/patterns.go` - All of the above plus agent-specific patterns: prompt injection, agent impersonation, tool abuse, trust boundary violations

Scanner patterns (353 regex patterns + ML): `aegisgate-scanner/pkg/detectors/` - Prompt injection: 19 | Secrets: 64 | PII: 85 | XSS: 12 | Dangerous code: 34 | Compliance: 65 | Path traversal: 19 | OT/ICS: 9 | RAG attack: 19 | Prompt leakage: 16

How to use them: Clone the repos, find the pattern files listed above, extract the regex strings, and create detection rules in your SIEM. Here are practical examples for the most common platforms:

Splunk — Credential Exfiltration to AI APIs

```
index=proxy_or_firewall
| regex _raw="AKIA[0-9A-Z]{16}"
| regex _raw="gh[pousr]_[A-Za-z0-9]{36}"
| regex _raw="sk-[A-Za-z0-9]{20}"
| regex _raw="-----BEGIN RSA PRIVATE KEY-----"
| regex _raw="xox[baprs]-[A-Za-z0-9-]+"
| table _time, src_ip, dest_host, _raw
```

> **Tune for your environment:** CI/CD pipelines (GitHub Actions, Jenkins) legitimately send API keys to AI APIs. Filter out known CI/CD source IPs — add `| where src_ip NOT IN ("10.0.1.50", "10.0.1.51")` for your build runner IPs, or `| where NOT match(src_host, "jenkins-*|runner-*|github-actions-*")` to exclude by hostname. Also filter QA/staging environments that legitimately use test API keys: `| where _raw NOT LIKE "%test_key%" AND _raw NOT LIKE "%dummy%"` . Start with block mode in monitor-only, review hits for a 24-hour window, then promote to alerting once false positive rate is acceptable.

Splunk — ARTEX Behavioral Indicators (Rapid API Reconnaissance)

```
index=proxy_or_firewall dest_host IN
("api.openai.com","api.anthropic.com","generativelanguage.googleapis.com")
| bucket _time span=1m
| stats count as api_calls, dc(dest_uri) as unique_endpoints by src_ip, _time
| where api_calls > 50 AND unique_endpoints > 10
| sort - api_calls
```

> **Tune for your environment:** Aggressive AI coding assistants can generate 40+ API calls/minute legitimately. Start with higher thresholds during business hours. The key signal is *diverse endpoint enumeration* — legitimate usage hits the same 2-3 endpoints; ARTEX probes many distinct paths.

Splunk — FortiBleed / Firewall Compromise Indicators

```
index=firewall_logs
| where src_ip IN
("193.8.186.7","80.75.212.113","213.21.239.65","208.94.246.58","69.195.129.144","96.45.42.1
73")
| table _time, src_ip, dest_ip, action, _raw
```

> **Tune for your environment:** These IOCs are from Kevin Beaumont's analysis. Rotate them as new IOCs emerge. Also hunt for Rclone egress to Wasabi/MEGA and RMM tools (MeshCentral, Atera, AnyDesk) installed on firewall management networks.

Splunk — Collabora Online Web Shell Hunting

```
index=web_logs uri="*proxy.php*"
| table _time, src_ip, uri, status, _raw
```

> Combine with a search for unexpected PHP files: `index=fileintegrity name="*.php" path="*nextcloud*" | where NOT match(name, "^(index|public|status|remote|ocs|cron|console)\.php$")`

ELK / OpenSearch — Credential Detection in AI Traffic

```
{
  "query": {
    "bool": {
      "filter": [
        {"terms": {"destination.domain": ["api.openai.com", "api.anthropic.com",
"generativelanguage.googleapis.com"]}}
      ],
      "must": {
        "regexp": {
          "request.body": "AKIA[0-9A-Z]{16}|gh[pousr]_[A-Za-z0-9]{36}|sk-[A-Za-z0-9]{20}"
        }
      },
      "must_not": [
        {"terms": {"source.ip": ["10.0.1.50", "10.0.1.51"]}}
      ]
    }
  }
}
```

Tune for your environment: The `must_not` filter excludes known CI/CD source IPs — populate the `source.ip` array with your build runner IPs. Add additional `must_not` entries for test environments: `{"match": {"request.body": "test_key"}}` and `{"match": {"request.body": "dummy"}}` to filter out QA traffic using placeholder credentials. Consider adding `"minimum_should_match": 1` to the `must` block if you combine multiple detection patterns. Filter by `event.category: "network"` to scope to proxy/egress logs only, excluding endpoint telemetry that may generate noise. If you run an egress proxy that decrypts TLS (e.g., mitmproxy, Squid with SSL bump), ensure the `request.body` field is populated — encrypted traffic will not match.

Microsoft Sentinel — Credential Exfiltration to AI APIs (KQL)

```
let ai_apis = dynamic(["api.openai.com", "api.anthropic.com",
"generativelanguage.googleapis.com"]);
let cicd_ips = dynamic(["10.0.1.50", "10.0.1.51"]);
CommonSecurityLog
| where DestinationHostName in~ (ai_apis)
| where SourceIP !in~ (cicd_ips)
| where RequestURL matches regex "(AKIA[0-9A-Z]{16}|gh[pousr]_[A-Za-z0-9]{36}|sk-[A-Za-z0-
9]{20})"
| project TimeGenerated, SourceIP, DestinationHostName, RequestURL, DeviceAction
| order by TimeGenerated desc
```

Tune for your environment: Populate the `cicd_ips` array with your actual CI/CD runner IPs — GitHub Actions, Jenkins, GitLab CI. Check the `DeviceAction` field: `"blocked"` entries from your firewall are already handled; focus on `"allowed"` connections where credentials reached the AI API. If you use Azure DevOps hosted agents, their IPs rotate — consider excluding by service tag `AzureDevOps` instead of static IPs. Also filter test environments: add `| where RequestURL !contains "test_key" and RequestURL !contains "dummy"` to reduce noise from QA traffic.

Sigma Rule — AI API Credential Exfiltration (Cross-Platform)

```
title: AI API Credential Exfiltration
id: 00000000-0000-0000-0000-000000000001
status: experimental
description: Detects AWS, GitHub, and OpenAI API keys being sent to AI API endpoints
references:
    - https://github.com/aegisgatesecurity/aegisgate-scanner
tags:
    - attack.exfiltration
    - attack.t1048
logsource:
    category: proxy
detection:
    selection_dest:
        destination.domain:
            - api.openai.com
            - api.anthropic.com
            - generativelanguage.googleapis.com
    selection_patterns:
        RequestURL|re:
            - 'AKIA[0-9A-Z]{16}'
            - 'gh[pousr]_[A-Za-z0-9]{36}'
            - 'sk-[A-Za-z0-9]{20}'
            - '-----BEGIN RSA PRIVATE KEY-----'
    condition: selection_dest and selection_patterns
falsepositives:
    - Legitimate API key rotation scripts (verify with dev team)
level: high
```

WAF Rules — Blocking Credential Exfiltration at the Edge

If your traffic flows through Cloudflare, AWS WAF, Akamai, or any WAF that supports custom regex rules, import the same patterns to block in real time rather than just detect:

```
# Cloudflare WAF custom rule (expression editor)
(http.request.body matches "AKIA[0-9A-Z]{16}") or
(http.request.body matches "gh[pousr]_[A-Za-z0-9]{36}") or
(http.request.body matches "sk-[A-Za-z0-9]{20}") or
(http.request.body matches "-----BEGIN RSA PRIVATE KEY-----")
```

> **Which patterns matter most:** The AWS key (`AKIA...`), GitHub token (`ghp_...`), and OpenAI key (`sk-...`) patterns are the highest-signal rules for credential exfiltration to AI APIs. The private key pattern catches broader exfiltration. Deploy in log-only mode first, review hits for 24 hours, then switch to block mode. Cloudflare Workers can also inspect response bodies — useful for detecting AI API responses that echo back credentials in error messages.

This gives you AegisGate's detection capability in the tools you already own, without deploying anything new. For the full 353-pattern set, clone the scanner repo and extract the regex strings from the detector files.

Step 4: Deploy Free AegisGate Tools (Option A)

Everything in Levels 1–4 below is **completely free, open source (Apache 2.0), with no feature gating or paid tiers**. Level 5 (Platform) has a free Community tier plus paid Pro+ tiers.

Not sure where to start? Use this table:

If you...	Start at	Time
Use ChatGPT/Claude/Gemini in a browser	Level 1	2 min
Use Copilot/Cursor/Claude Code in an IDE	Level 1 + 2	12 min
Build AI apps with RAG pipelines or agents	Level 1 + 2 + 3	27 min
Run MCP servers for AI agents	Level 4	20 min
Run AI coding agents (Codex, Claude Code)	Level 3 (scan instruction files)	5 min
Need org-wide AI traffic visibility	Level 5	Hours
Aren't sure what any of this means	Level 1 (it's a browser extension)	2 min

---

Level 1: Browser Extension — For Everyone (2 minutes)

Tool: AegisGate Lens

If you're using ChatGPT, Claude, Gemini, or any AI chat in your browser — install Lens. It's a browser extension that runs **100% on-device**. Zero data leaves your browser. No telemetry by default. 219 regex patterns + ML detection.

How to install: - **Chrome**: Chrome Web Store — requires Chrome 116+ - **Firefox**: Firefox Add-ons — requires Firefox 142+

Level 2: Local Proxy — For Developers Using AI Coding Tools (10 minutes)

Tool: AegisGate Rampart

If your developers use GitHub Copilot, Cursor, Claude Code, or local LLMs — Rampart sits between their editor and the AI model. 302 regex patterns + ML detection. Monitor mode or Block mode. Zero telemetry. Runs fully air-gapped.

> **For the IT generalist — what is a MITM proxy?** MITM stands for "man in the middle." In this context, it's a local proxy — a small program running on the developer's machine that intercepts traffic between their editor and the AI model. Think of it as a security checkpoint: every request the editor sends to OpenAI or Anthropic passes through Rampart first. Rampart scans it for secrets, PII, prompt injection, and other risks. In monitor mode, it logs alerts. In block mode, it stops the request before it reaches the AI server. The developer's workflow doesn't change — Rampart is transparent to the editor.

```
# Download pre-built binary from GitHub Releases
# https://github.com/aegisgatesecurity/aegisgate-rampart/releases/tag/v0.9.3

# Or build from source
git clone https://github.com/aegisgatesecurity/aegisgate-rampart
cd aegisgate-rampart
go build -o rampart ./cmd/rampart
./rampart --mode monitor
```

Level 3: Static Analysis — For CI/CD Pipelines (15 minutes)

Tool: AegisGate Scanner

353 regex patterns + ML. Scans AI artifacts for prompt injection, secrets, PII, XSS, dangerous code, compliance violations, path traversal, OT/ICS, RAG attacks, and prompt leakage. Outputs SARIF, SBOM, JSON, and 5 other formats.

> **For AI developers — RAG pipeline poisoning:** If your AI application uses RAG (Retrieval-Augmented Generation), the documents in your vector database are an attack surface. An attacker who can inject content into your knowledge base — through a compromised document, a poisoned web scrape, or a supply chain attack on a data source — can embed prompt injection payloads that execute when users query the system. Scanner's 19 RAG attack patterns detect injected instructions in retrieved chunks before they reach the model. Scan all documents before ingestion, and scan retrieved chunks at query time. The same applies to A2A (agent-to-agent) communication — agents receiving instructions from other agents should validate those instructions the same way you validate user input.

```
# One-line installer
curl -fsSL https://raw.githubusercontent.com/aegisgatesecurity/aegisgate-
scanner/main/install.sh | bash

# Scan a directory
./scanner scan ./ai-artifacts/ --format sarif --output results.sarif

# GitHub Action
# uses: aegisgatesecurity/aegisgate-scanner@v0.4.3
```

Level 4: Secure MCP Servers — For Teams Running MCP-Based Agents (20 minutes)

Tool: AegisGate MCP

Hardened MCP server framework with 23 security layers. 79 regex patterns + ML. Zero external dependencies (pure Go stdlib). Distroless Docker image with zero CVEs. Runs fully air-gapped.

Level 5: Enterprise Gateway — For Organizations Running AI at Scale (hours)

Tool: AegisGate Platform

Self-hosted gateway. Every request and response scanned across 6 protection pillars: HTTP, MCP, A2A, ACP, Response, Trust.

Free Community tier: 303 regex patterns, protocol scanning, 4 compliance frameworks, fail-closed by default, zero external dependencies, air-gapped capable.

Pro+ (paid): ML anomaly detection, SIEM integration, additional compliance frameworks, HSM/FIPS, Kubernetes clustering.

Performance: 23,578 RPS with 0% errors. -2.8ms p99 overhead. Blocked requests resolve in 7.2ms p50.

```
git clone https://github.com/aegisgatesecurity/aegisgate-platform
cd aegisgate-platform
go build -o aegisgate-platform ./cmd/aegisgate-platform/
./aegisgate-platform setup --non-interactive
./aegisgate-platform --config aegisgate-platform.yaml --embedded-mcp
```

> **Air-gap note:** All AegisGate tools run fully air-gapped. All detection is local. No cloud, no phone-home, no telemetry by default. This is not a feature — it's a requirement for OT, defense, government, and regulated environments.

How to Stop the Markdown-as-Executable Attack Today

This is a concrete, actionable response to the agent instruction file poisoning attack disclosed Oct 22. No patch exists for the underlying issue — agents are designed to follow `AGENTS.md` and `CLAUDE.md` files. But you can detect and block the malicious instructions before agents read them.

What to Do Right Now

1. Check your Downloads folder:

```
find ~/Downloads -name "AGENTS.md" -o -name "CLAUDE.md" -o -name ".codex" -o -name
".claude" 2>/dev/null
```

2. Scan with AegisGate Scanner before starting agent sessions:

```
./scanner scan ~/Downloads/ --format json --output scan-results.json
```

3. Add a GitHub Action to scan repos for malicious instruction files:

```
# .github/workflows/agent-instruction-scan.yml
name: Agent Instruction File Security Scan
on: [push, pull_request]
jobs:
  scan:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: aegisgatesecurity/aegisgate-scanner@v0.4.3
        with:
          path: ./
          format: sarif
```

4. Train your team on the IOC: - `function-color` npm package (removed from npm, copies may exist) - Any `AGENTS.md` instructing Codex to modify `~/.codex/config.toml` is suspicious - Any `CLAUDE.md` instructing Claude to change API endpoints or install packages is suspicious - If an agent suddenly tries to install packages or modify config files, stop and investigate

5. Stop running agents in Downloads: - Create dedicated project directories - Configure Chrome/Edge to prompt before downloading

---

Step 5: ARTEX-Specific Guidance — What to Look For

If you're hunting for ARTEX-like activity in your environment, here's what we know from Zscaler ThreatLabz and Microsoft threat intelligence:

Behavioral Indicators

1. **Rapid reconnaissance**: Many API calls in a short window, probing multiple endpoints. Storm-0940 hit 2,900 auth attempts/minute.
2. **Credential attacks through APIs**: Password spraying, credential stuffing via application credentials. 151,000+ failed sign-in attempts.
3. **AI-accelerated activity**: Attack progression faster than human-operated. Storm-0940: initial targeting to privileged access in under 3 hours.
4. **Multi-agent patterns**: Coordination between agents — one doing recon, another exploitation, a third data access. GLM reviewed DeepSeek's subagent work.
5. **Microsoft Graph anomalies**: Reconnaissance queries, application credential changes, new OAuth app registrations.
6. **Cross-AI-session handoffs**: Custom scripts sharing data between different AI tools.

Technical Indicators

- Unexpected Go binaries on workstations (ARTEX is Go + Next.js)
- API calls to OpenAI-compatible endpoints from unexpected sources
- MCP traffic (Streamable HTTP/SSE) to external servers
- Configurations specifying DeepSeek, GLM, or Grok models
- Scripts coordinating handoffs between AI sessions
- 878+ ARTEX forks exist — check developers' machines

Establish Your Baseline First

Before hunting for anomalies, establish what normal AI API traffic looks like in your environment. Which hosts legitimately call AI APIs? What's your baseline volume and timing? Anything outside that baseline since October 17 is your hunting ground.

What ARTEX Means for You

ARTEX is open source. It's been forked 878+ times. The original repo was deleted, but the code is everywhere. Multiple threat groups are using it, not just Storm-0940. Microsoft observed attackers from "several geographic regions" using the framework.

The implication: AI-accelerated attacks are no longer the exclusive domain of nation-states. The tooling is freely available. Expect to see it used by cybercrime groups, ransomware operators, and opportunistic attackers.

This is why machine-speed attacks require machine-speed detection — it's not a marketing tagline, it's a survival requirement. Human-paced incident response cannot keep up with AI agents that go from reconnaissance to privileged access in under 3 hours.

Your detection needs to be automated, not manual. Your response needs to be scripted, not ad-hoc. And your AI traffic needs to be visible, not invisible.

Step 6: Sector-Specific Guidance

The threats above affect everyone, but the compliance triggers, log sources, and priorities differ by sector.

OT/ICS & Industrial

ARTEX automates the full attack chain. The same multi-agent approach that hit 2,250 M365 tenants could target SCADA, DCS, and ICS environments.

What to do: - If AI touches your OT environment — even indirectly — deploy Rampart on those developer machines. Air-gapped, zero network dependency. - Route AI agents through Platform. 303 regex patterns include Modbus, DNP3, OPC-UA detection. - Scan AI training data and agent skills with Scanner before deployment. - Install Lens on HMI workstations.

Defense & Government

Storm-0940 is a Chinese APT. Defense contractors and federal agencies are prime targets. NIST 800-171 3.13.1 requires network isolation for CUI environments. SaaS AI security tools violate this.

What to do: - Deploy Platform air-gapped within your security boundary. Community tier is free. - MCP's distroless Docker image aligns with STIG-style hardening. - Rampart on developer workstations with zero network dependency. - Rotate ConnectWise credentials. **Verify Secure Boot DBX revocations** (see BlackX sidebar).

Healthcare & HIPAA

If a clinician pastes a patient case into ChatGPT, that's a HIPAA reportable event. Lens catches 77 PII patterns before it hits send.

What to do: - Deploy Lens on all clinical and admin workstations (2 minutes each). - Scan RAG documents with Scanner before ingestion. - **Retroactive exposure check**: Search proxy logs for AI API traffic from PHI-handling workstations since October 17. - **If you use Fortinet or Palo Alto VPN for remote clinical access**: Rotate all VPN credentials immediately (FortiBleed campaign).

Financial Services

AegisGate detects SWIFT/BIC codes, credit card numbers, and bank accounts in AI traffic.

What to do: - Deploy Lens on trader, analyst, and customer service workstations. - Deploy Rampart on developer machines. - **Retroactive exposure check**: Search proxy logs for card numbers and SWIFT codes in AI API traffic. - **If you use Fortinet or Palo Alto**: Rotate all VPN credentials. Check for dormant admin accounts. The k0t campaign is linked to Qilin ransomware — verify your firewall integrity.

Cloud & Multi-Cloud

AI API calls from your cloud infrastructure are probably invisible to you.

What to do: - **AWS**: Search CloudTrail for `bedrock:*`, `sagemaker:*`, `InvokeModel`. Check VPC Flow Logs for AI API endpoints. - **Azure**: Check Entra ID sign-in logs for Storm-0940 indicators. Search Graph audit logs. - **GCP**: Search Cloud Audit Logs for `aiplatform.*` and Vertex AI calls. - Deploy Platform as a cloud gateway to intercept and scan AI traffic at the infrastructure level.

State, Local & Education

Universities and local governments run PHP extensively and are heavy AI users. Both sectors are often resource-constrained. ~30,800 eMudhra-affected certificates are 99.9% via InCommon/Internet2 — higher-ed is directly impacted.

What to do: - **Patch PHP immediately.** PHP 8.5.12, 8.4.27, 8.3.36, or 8.2.35. - **Reissue InCommon certificates before Oct 27** (Chrome drops eMudhra trust). - Install Lens on all faculty, staff, and admin workstations. - Rotate ConnectWise credentials if applicable. - Deploy Platform Community tier (free) for AI traffic visibility. - **If you run Fortinet or Palo Alto firewalls**: Rotate VPN credentials, hunt for dormant admin accounts.

Quick Reference: What to Do Right Now

If you have 5 minutes:

1. Install AegisGate Lens browser extension — Chrome or Firefox
2. Check if your Node.js and Bun installations are patched for eMudhra
3. Verify your PHP version — if it's not 8.5.12, 8.4.27, 8.3.36, or 8.2.35, start planning downtime NOW
4. Check your GitLab version — if it's not patched, restrict network access immediately
5. **Check your Downloads folder for `AGENTS.md` or `CLAUDE.md` files you didn't create**
6. **If you run Fortinet or Palo Alto firewalls: start rotating VPN credentials NOW**

If you have 30 minutes:

1. All of the above
2. Download AegisGate Rampart for your dev team
3. Install AegisGate Scanner and scan Downloads + working directories
4. Pull AegisGate detection patterns and import into your SIEM (see Step 3)
5. Audit Microsoft Graph logs for Storm-0940 indicators
6. **Upgrade Collabora Online richdocumentscode to 26.4.303 if you run Nextcloud**
7. **Patch Arista VeloCloud if you run on-prem VCO** (CISA KEV due Oct 26)
8. **Hunt for dormant admin accounts on Fortinet/Palo Alto firewalls**
9. **Verify Secure Boot DBX revocations on Windows systems** (see BlackX sidebar)

If you have a few hours:

1. All of the above
2. Deploy AegisGate Scanner in your CI/CD
3. Add the Scanner GitHub Action to scan repos for malicious instruction files in PRs
4. Evaluate AegisGate MCP for your MCP infrastructure
5. Stand up AegisGate Platform Community tier as a monitoring gateway
6. **Run Nextcloud integrity scan and hunt for web shells** if you use Collabora Online
7. **Check owned.lab6.com for your organization's IPs** if you run Fortinet/Palo Alto

If you're in active incident response:

1. Check PHP access logs back to **October 17** for PHP2Shell exploitation indicators
2. Revoke and rotate **all** ConnectWise-related credentials. Check for OAuth persistence
3. Scan LiteLLM installations for compromise (unexpected `.pth` files, outbound connections to `litellm[.]dev`)
4. Hunt for ARTEX behavioral indicators in API logs (see Step 5)
5. **Hunt for FortiBleed indicators**: IOC IPs in firewall logs, RMM tools on management networks, Rclone egress to Wasabi/MEGA, dormant admin accounts
6. **Hunt for Collabora web shells**: Nextcloud integrity scan, unexpected PHP files, proxy.php in web logs during CODE server downtime
7. **Check VCO logs for SQLi** from `74.82.197.86`, `38.54.50.151` if you run Arista VeloCloud

8. Check all developer machines for OpenCode versions < 1.18.33
9. Scan all Downloads folders for planted `AGENTS.md` / `CLAUDE.md` files
10. Import AegisGate regex patterns into SIEM and run retroactive searches
11. Establish your AI traffic baseline and flag anything outside it since October 17

Appendix: Log Sources for Hunting

Incident	Log Source	What to Search For
PHP2Shell	nginx/apache access logs, PHP-FPM logs, web server error logs	Unusual `multipart/form-data` POST requests, unexpected process spawning from web workers, outbound connections to unfamiliar IPs. **Check back to Oct 17.**
Storm-0940	Entra ID sign-in logs, Microsoft Graph audit logs, M365 Unified Audit Log	Rapid auth bursts (2,900+/min), application credential changes, new OAuth app registrations, privilege modifications
GitLab RCE	`/var/log/gitlab/gitlab-rails/production.log`, `/var/log/gitlab/nginx/gitlab_access.log`	Unauthenticated requests, error spikes, unexpected process execution
LiteLLM	Process network connections, Python `site-packages` directories, container egress logs	`cook.pth`, `litellm_init.pth`, `sitecustomize.py` (unexpected), outbound to `litellm[.]dev`, GitHub token usage Oct 10-21
AI credential exfil	Egress proxy logs, firewall logs, CloudTrail, Azure Monitor, GCP Cloud Audit Logs	API keys (`AKIA...`, `ghp_...`, `sk-...`) in traffic to AI API endpoints
ARTEX	API gateway logs, MCP server logs, endpoint detection logs, process inventory	Rapid multi-endpoint reconnaissance, Go binaries in unexpected locations, MCP connections to external servers
Arista VeloCloud	VCO web access logs	SQLi patterns (union-based, time-based blind) from `74.82.197.86`, `38.54.50.151`
Collabora Online	Nextcloud integrity scan results, web server access logs	`EXTRA_FILE` entries in integrity check, unexpected PHP files (e.g., `Nx_33143.php`), requests to `proxy.php` during CODE server downtime
Enterprise Firewall/VPN	FortiGate System→Events (filter "config"), Palo Alto logs, firewall management network logs	IOC IPs, dormant admin accounts, RMM tools (MeshCentral, Atera, AnyDesk, Teramind), AdaptixC2, Rclone egress to Wasabi/MEGA

Incident	Log Source	What to Search For
BlackX UEFI	UEFI/Secure Boot configuration, boot logs, EDR process logs	Missing DBX revocations for CVE-2022-21894, `inst.exe` in `C:\ProgramData\Intel\`, `archivetj[.]com` C2, sideloaded `googleupdate.exe`
Markdown-as-Executable	File integrity monitoring, Downloads folder monitoring, process logs	New `AGENTS.md` or `CLAUDE.md` in Downloads, config file modifications, `function-color` npm package
OpenCode MCP RCE	Process logs, PowerShell event logs (Event ID 4104), network connections	PowerShell execution from OpenCode process, unexpected browser launches during `opencode mcp auth`

Appendix: Indicators of Compromise (IOCs)

Known indicators associated with Black October incidents. Use these alongside the SIEM queries in Step 3 for retroactive hunting.

> **Disclaimer**: IOCs are point-in-time. Attackers rotate infrastructure. Use these for retroactive hunting (Oct 17 onward), not as your only detection layer.

PHP2Shell

Type	Indicator	Notes
Network	Unusual `multipart/form-data` POST requests to PHP applications	Primary attack vector — check back to Oct 17
Process	Unexpected process spawning from web worker (www-data, php-fpm)	Shell execution from PHP runtime
Network	Outbound connections to unfamiliar IPs from web server	Post-exploitation C2 or data exfil
Config	`session.upload_progress.enabled` still on	Attack path — disable even after patching (bypass exists)
Config	`session.serialize_handler` = `php_binary` / `php_serialize` then switched at runtime	CVE-2026-115560 requirement — unusual config

eMudhra CA Compromise

Type	Indicator	Notes
Certificate	TLS certs issued by `emSign ECC Root CA - G3` intermediates	Chrome drops trust Oct 27 — replace before then

Type	Indicator	Notes
Certificate	Any cert chain through `eMudhra Root CA - G3`	30,800 affected certs, 99.9% via InCommon/Internet2
Network	Successful TLS connections that should fail (forged certs)	If trust store still trusts eMudhra, forged certs validate
IoT	Fire TV Stick, unpatched RHEL/Mint accepting forged certs	Dormann MITM testing — IoT/embedded devices are active targets

Arista VeloCloud Orchestrator

Type	Indicator	Notes
Network	`74.82.197.86`	Attacker IP — actively exploited
Network	`38.54.50.151`	Attacker IP — actively exploited
Log	VCO web access logs — SQLi patterns	Union-based injection, time-based blind
Behavioral	Privilege escalation to superuser via SQLi	Authenticated user → superuser

Collabora Online CODE Server

Type	Indicator	Notes
File	Unexpected PHP files in Nextcloud data/app directory	Web shells dropped via proxy.php RCE
File	`Nx_33143.php`	Example web shell found via integrity scan
Log	Requests to `proxy.php` when CODE server not running	Exploitation window — during restart/upgrade/crash
App	richdocumentscode < 26.4.303	Vulnerable version — must upgrade or remove (disabling insufficient)

Enterprise Firewall/VPN (k0t/FortiBleed)

Type	Indicator	Notes
Network	`193.8.186.7`	Attacker C2 IP (Kevin Beaumont)
Network	`80.75.212.113`	Attacker C2 IP
Network	`213.21.239.65`	Attacker C2 IP
Network	`208.94.246.58`	Attacker C2 IP
Network	`69.195.129.144`	Attacker C2 IP
Network	`96.45.42.173`	Fortinet's own ASN (co-opted)

Type	Indicator	Notes
Log	FortiGate System→Events, filter "config"	Configuration changes indicating compromise
Process	MeshCentral, Atera, AnyDesk, Teramind	RMM tools used for persistence
Process	AdaptixC2	C2 framework
Network	Rclone egress to Wasabi, MEGA	Data exfiltration
Account	Dormant admin accounts	Backdoor persistence — "staggering amount" affected
External	`owned.lab6.com`	Victim IP list — check if your IPs are listed

LiteLLM Supply Chain Attack

Type	Indicator	Notes
C2 Domain	`litellm[.]dev/healthcheck`	Exfiltration endpoint — HTTP POST with stolen credentials
File	`cook.pth` in site-packages	Imports `cook.py` on Python startup
File	`cook.py` in site-packages	`privcollect` stealer — system, browser, dev, cloud, AI secrets
File	`litellm_init.pth`	Earlier variant (June 2026)
File	`sitecustomize.py` (unexpected)	C2 dead drop check-in every hour (Linux/macOS)
Network	`69.48.229.143`	Attacker IP (root cause: unrevoked master key)
Repo	`Berri-AI/litellm` (with hyphen)	Lookalike repo used to hijack release pipeline
Behavioral	Short-lived branches `ci-tmp-a351cc`, `ci-tmp-4f8517`	Malicious GitHub Actions workflow branches
Cloud	AWS `GetCallerIdentity` from unknown IPs	Stolen credentials validated within 4.5 hours

Storm-0940 / ARTEX

Type	Indicator	Notes
Behavioral	2,900+ auth attempts/minute from single source	Password spraying against M365 tenants
Behavioral	Privileged cloud access in <3 hours from initial targeting	AI-accelerated attack timeline
Process	Unexpected Go binaries in temp dirs or user profiles	ARTEX is written in Go + Next.js
Network	API calls to OpenAI-compatible endpoints from workstations	ARTEX uses DeepSeek, GLM, Grok via OpenAI-compatible gateway

Type	Indicator	Notes
Network	MCP traffic (Streamable HTTP/SSE) to external servers	ARTEX supports MCP
OAuth	New OAuth app registrations in M365 after Oct 17	Persistence mechanism

BlackX UEFI Bootkit

Type	Indicator	Notes
File	`inst.exe` in `C:\ProgramData\Intel\`	Bootkit installer
File	`PVMDIJGKAX2`	Encrypted payload (extensionless)
Crypto	RC4 key: `Q46F6MDLI3RX2KV1`	Decryption key for payload
Network	`archivetj[.]com`	C2 domain for Rust payload plugins
Process	`googleupdate.exe` (sideloaded) + `OxiDLLRunner.dll`	DLL sideloading via legitimate binary
Network	Google Drive, OneDrive, Yandex.Disk	C2 channels blending with enterprise traffic
Config	Missing Secure Boot DBX revocations for CVE-2022-21894	Systems "patched" but without manual DBX update — still vulnerable

Markdown-as-Executable

Type	Indicator	Notes
File	`AGENTS.md` in `~/Downloads`	Planted instruction file — agents read and obey
File	`CLAUDE.md` in `~/Downloads`	Planted instruction file — agents read and obey
npm	`function-color`	Malicious npm package (removed)
Config	`~/.codex/config.toml` modifications	Malicious MCP server entries
Config	`.claude/settings.json` modifications	API endpoint changes redirecting traffic

OpenCode MCP RCE

Type	Indicator	Notes
Process	PowerShell execution from OpenCode process	Event ID 4104 — command injection
Version	OpenCode < 1.18.33	Vulnerable (1.0.137–1.18.32) — no CVE
Cloud	API keys in OpenCode Console	1.4% of hosted workspaces — rotate keys

---

What AegisGate Cannot Do (Honest Limits)

I'm not going to pretend one tool solves Black October. Here's what AegisGate does NOT do:

- **Patch your PHP servers.** You have to do that. PHP 8.5.12, 8.4.27, 8.3.36, or 8.2.35. Now. And keep `session.upload_progress.enabled=0`.
- **Manage your TLS certificate trust stores.** Update your CA bundles. Check Node.js and Bun. Test against Mozilla's incident endpoints. Check IoT devices.
- **Patch your GitLab.** Update to 19.5.1, 19.4.3, 19.3.6, 19.2.9, 19.1.9, 19.0.10, or 18.11.13.
- **Patch your Arista VeloCloud.** VCO 5.2.3.18+, 6.1.3.10+, 6.4.2.12+.
- **Upgrade your Collabora Online.** richdocumentscode 26.4.303. Then hunt for web shells.
- **Rotate your firewall VPN credentials.** That's on you. Check for dormant admin accounts.
- **Verify your Secure Boot DBX revocations.** Microsoft's manual revocation steps are on you.
- **Fix Microsoft's SharePoint patches.** That's on Microsoft. Restrict access and wait.
- **Stop a website from downloading a file to your Downloads folder.** Browser download behavior is controlled by the browser. Scanner CAN detect the malicious file before an agent reads it.
- **Replace your SIEM.** AegisGate is a detection and prevention layer for AI agent traffic. In the free tiers, you can extract our patterns and feed them into your SIEM manually.
- **Stop ARTEX running on someone else's machine.** AegisGate detects the *effects* of ARTEX in your traffic. It can't disable the attacker's tooling.
- **Detect novel AI attack patterns without ML (free tier).** The 303 regex patterns in Platform's Community tier catch known patterns. The ML anomaly detection is a Pro+ feature. Lens and Rampart include ML detection for free.

What AegisGate DOES do: it gives you visibility into a layer of traffic that you almost certainly have zero visibility into right now — the AI agent layer. And Black October just proved that layer is being actively exploited.

About This Guide

This guide was written by a blue teamer who built AegisGate because they saw the AI agent blind spot coming. Everything in Levels 1–4 is free and open source (Apache 2.0) with no paid tiers, no feature gating, and no "contact sales for the real version." Level 5 (Platform) has a free Community tier that includes 303 regex patterns and protocol-level scanning — the ML anomaly detection and SIEM integration are Pro+ features, and I've been upfront about that.

There is no upsell here. There is no "schedule a demo." There are repos you can clone and tools you can deploy in minutes — including in air-gapped, regulated, and classified-adjacent environments where SaaS tools can't go.

The repos: - **Lens** (browser extension, 219 regex + ML): github.com/aegisgatesecurity/aegisgate-lens - **Rampart** (local proxy, 302 regex + ML): github.com/aegisgatesecurity/aegisgate-rampart - **Scanner** (static analysis, 353 regex + ML): github.com/aegisgatesecurity/aegisgate-scanner - **MCP** (secure MCP framework, 79 regex + ML): github.com/aegisgatesecurity/aegisgate-mcp - **Platform** (enterprise gateway, 303 regex Community / +ML Pro+): github.com/aegisgatesecurity/aegisgate-platform

If this helps you, star the repos so other blue teamers find them. If it doesn't help you, the detection patterns are still free — take them and use them in your SIEM. No strings.

Stay safe out there.

This guide will be updated as the situation develops. Last updated: October 24, 2026 06:20 UTC.